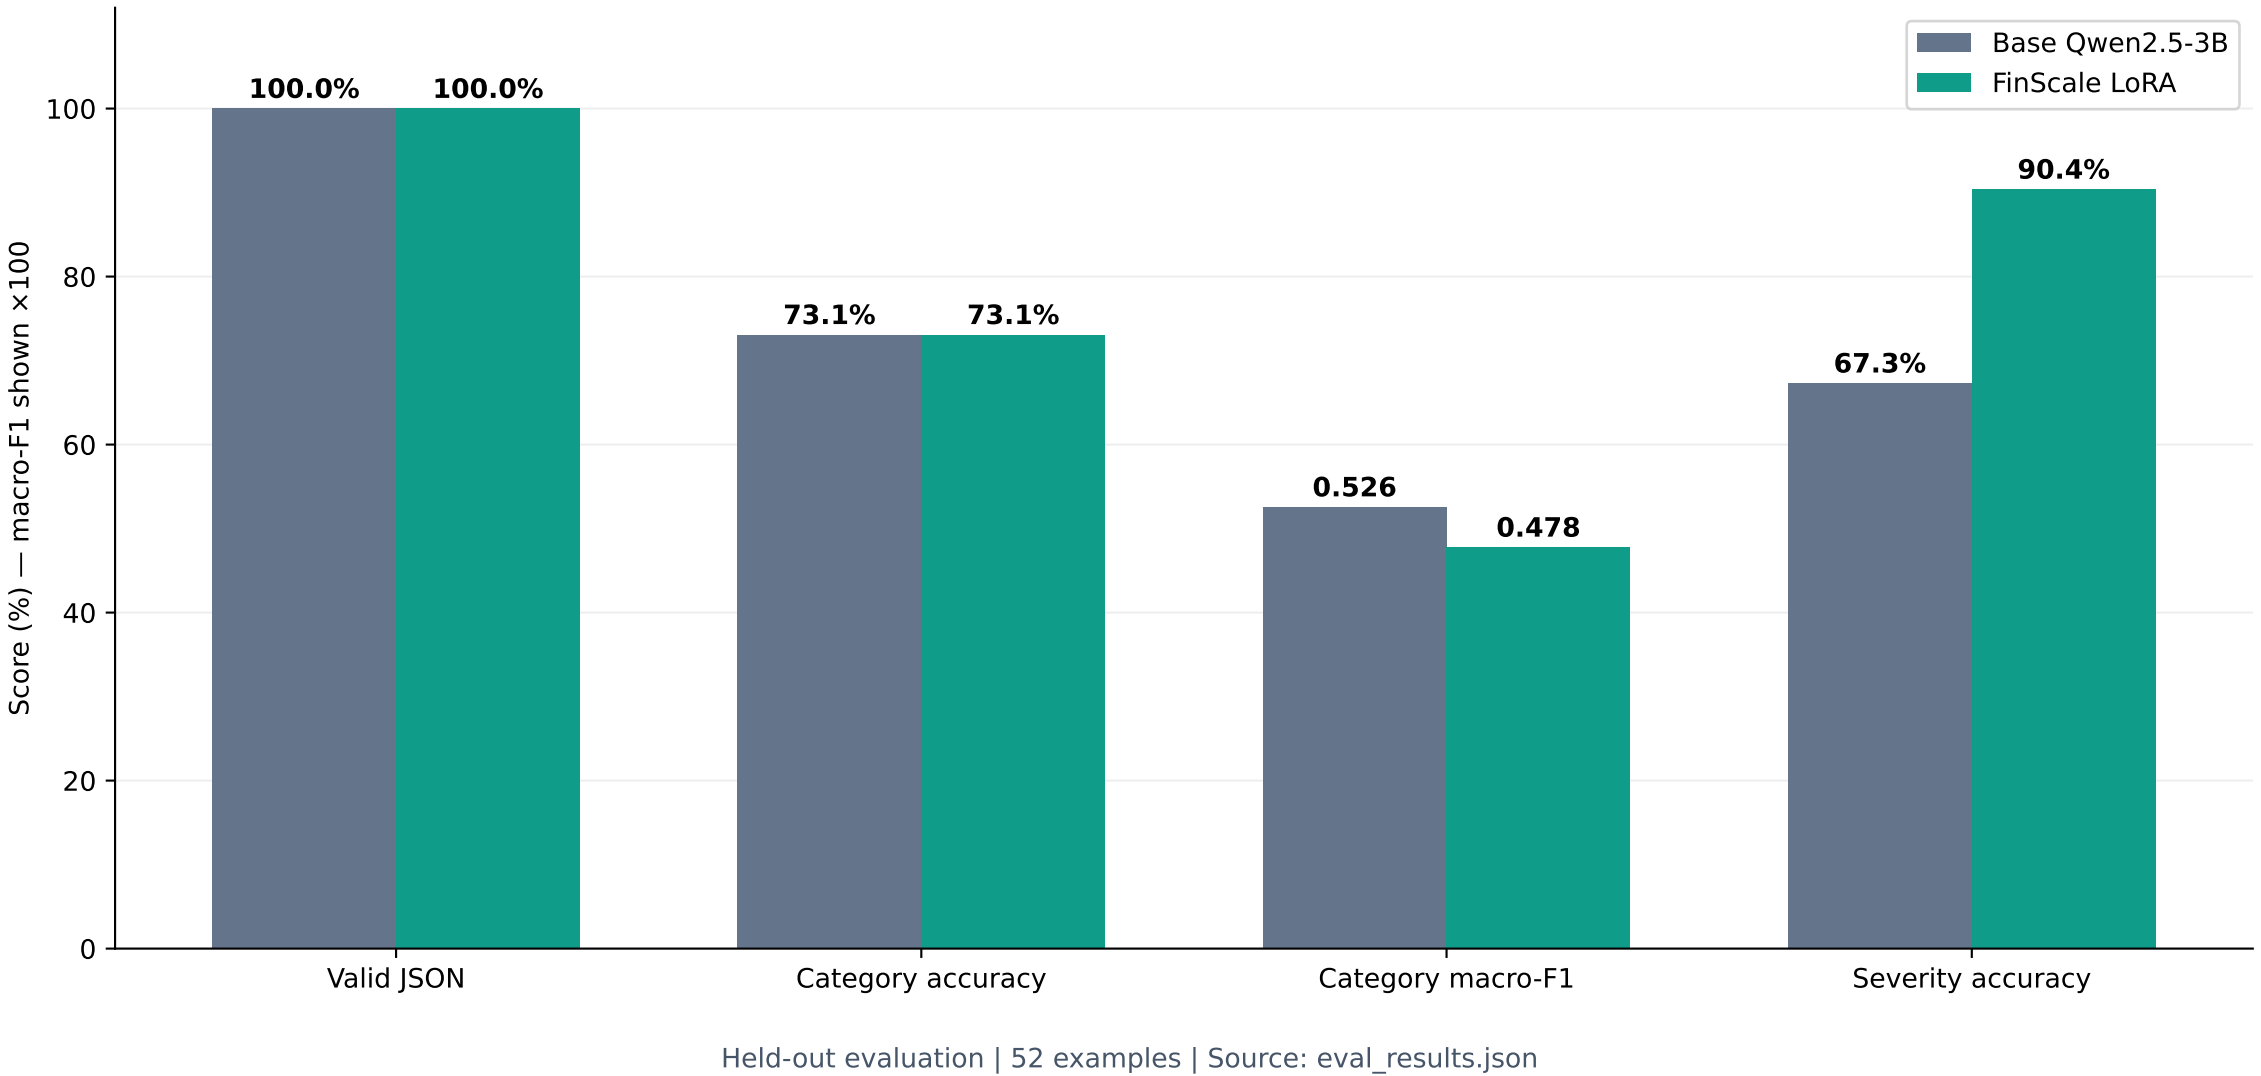
FinScale | Held-out financial risk evaluation
Base vs LoRA Fine-Tuned Qwen2.5-3B
Base Qwen2.5-3B
FinScale LoRA
Score (%) — macro-F1 shown ×100
0
20
40
60
80
100
100.0%
100.0%
73.1%
73.1%
0.526
0.478
67.3%
90.4%
Valid JSON
Category accuracy
Category macro-F1
Severity accuracy
Held-out evaluation | 52 examples | Source: eval_results.json